Varma Mutual Pension Insurance Company

# Three models in production on one reusable MLOps pattern, built with your new specialists

Draft staffing revision - candidate evidence, scope and rates require owner confirmation

# Three MLOps use cases over six months, with a funded three-role team

- **One MLOps golden path** on your existing cloud data platform replaces isolated deployment scripts and per-model configurations
- **Quick win:** use case 1 in production with monitoring by the end of month 2
- **Your new specialists own it:** they co-build use case 2 and build use case 3 themselves, with us reviewing
- **Three complementary roles, subject to candidate verification;** your right to require replacement accepted
- **Draft EUR 199,350 capacity fee**, EUR 66,450 per two-month milestone; scope, data readiness and risk allowance require confirmation.
- Leadership quick win: lead time from approved model to production, measured from use case 1 and reported monthly

# One MLOps golden path for every model, instead of one-off deployment work per experiment

- Every model, current and future, uses the same repo template, configuration schema, pipeline and monitoring. A new use case starts from the template, not from a blank script.
- Built on your existing cloud data platform; MLflow-compatible registry, so the pattern is not tied to one vendor.

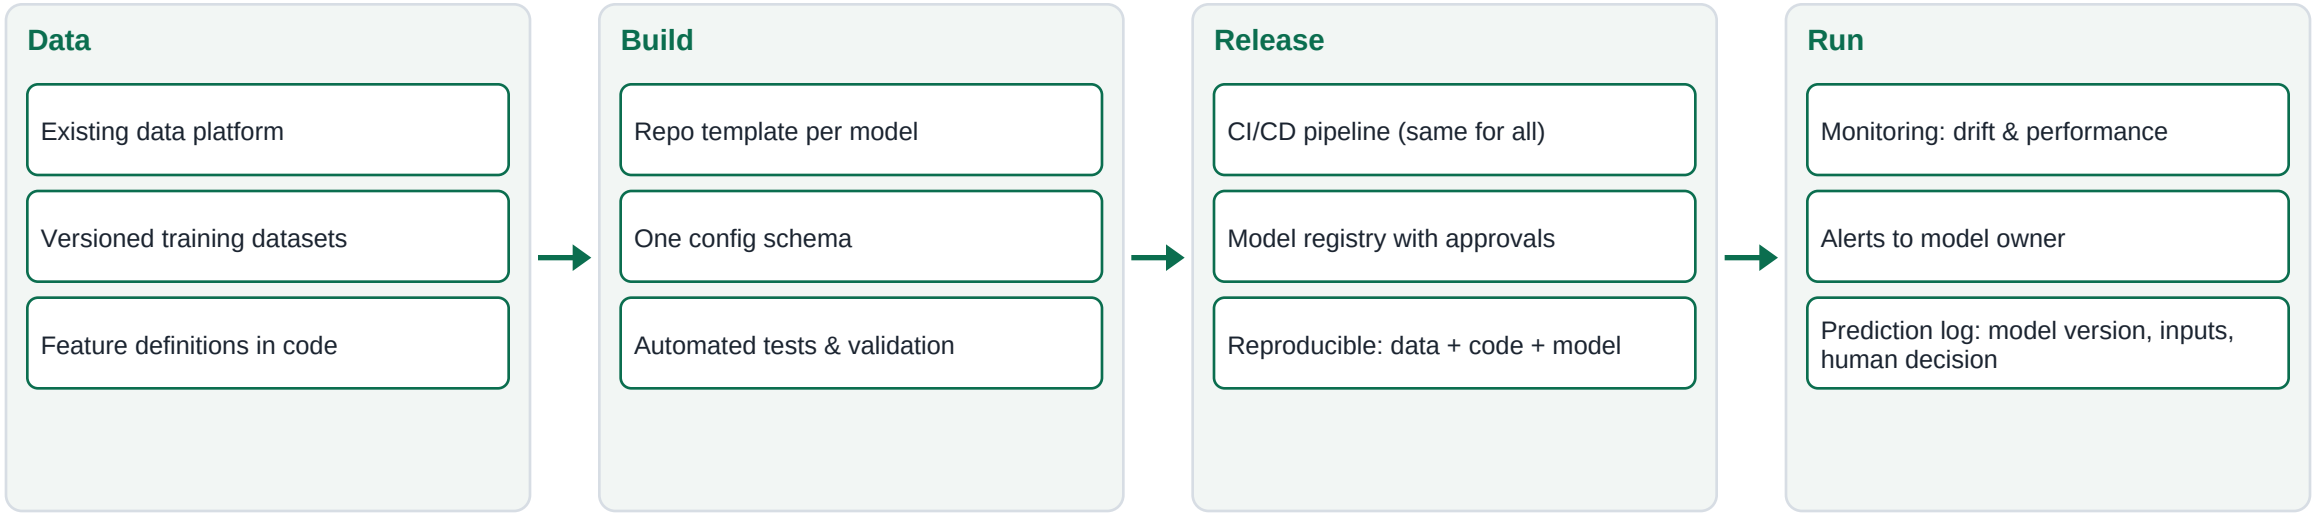


Models support pension decisions, they do not make them: shadow run first, then a human-approved release; rollback to the previous registered model in one step

# Use case 1 in production by month 2; your new specialists ship use case 3 themselves

| Month | Use case | Who builds | Technical proof |
|---|---|---|---|
| 1–2 | Use case 1 on the golden path (quick win) | Recordly, your specialists pairing | Live with drift alerts; performance checked when outcomes arrive (labels lag) |
| 3–4 | Use case 2 on the same template | Your specialists, Recordly pairing | Built without new deployment scripts |
| 5–6 | Use case 3 + triage of the other ~17 experiments by data readiness, value and risk class | Your specialists lead, Recordly reviews | Handover check: your team releases unaided |

Gate 1 (week 3): use case 1 has usable data, a risk class and a named model owner; otherwise we swap to the next use case. Technical unknowns we price in: which of the three use cases has production-ready data; access to personal data for training (minimised, pseudonymised); EU AI Act classification per use case. Documentation uses Finnish pension terms (työeläke, TyEL, eläkepäätös), no translation needed.

# Proposed role mix, allocations and verification before commitment

- **Replacement:** at your request, equal seniority within 10 working days, 2-week overlap at our cost
- **Delivery governance:** monthly steering with a written status (use cases, model health, risks); no weekly status meetings
- **Handover:** runbooks, the template and onboarding guide in Finnish, in your repos
- You may recruit our consultants after 12 months, with no fee

| Proposed role | Role | Allocation | Required certificates | Target start |
|---|---|---|---|---|
| Lead ML engineer | Lead ML engineer | 100% | Databricks Certified Machine Learning Professional; Microsoft Azure Data Scientist Associate | At signing |
| MLOps / platform engineer | MLOps / platform engineer | 60% | Databricks Data Engineer Professional; Microsoft DevOps Engineer Expert | At signing |
| Solution architect | Solution architect, delivery governance | 15% | Microsoft Azure Solutions Architect Expert | At signing |

Draft role profiles: nominate consultants and verify CVs, certificates, languages, dates and actual availability before submission. No person is booked by this draft.

# EUR 199,350 with all three AI delivery roles priced

- 20 working days/planning month; ML 100%, platform DE 60%, architect 15%. Confirm exact calendar and availability.
- Proposed ML and architect base EUR 125/h + EUR 5/h AI = EUR 130/h; DE EUR 115/h + EUR 5/h AI = EUR 120/h. Non-DE bases require owner confirmation.
- Three use cases prioritised at discovery; replace blocked use cases within the capacity limit. Scope growth and risk allowance require written agreement.

| Role | Six-month effort / rate | Fee |
|---|---|---|
| Lead ML engineer | 120 days x 7.5 h x EUR 130 | EUR 117,000 |
| MLOps / platform DE | 72 days x 7.5 h x EUR 120 | EUR 64,800 |
| AI solution architect | 18 days x 7.5 h x EUR 130 | EUR 17,550 |
| Total excl. VAT | 210 days; 3 x EUR 66,450 | EUR 199,350 |

No new licences assumed, subject to client confirmation. Candidate credentials and availability pending; VAT, cloud and travel excluded. No promise to productionise all 20 experiments.

# Staffing review and decision before submission

- Rauno / Codex staffing review: maximum two architects respected; complementary roles included; DE/DA floors and AI premium applied to the stated role scope.
- This revision replaces the lower draft price for the same capacity. If the buyer budget is lower, reduce scoped work or hours and re-estimate; do not discount below hourly minima.
- Candidate CVs, credentials, language skills, exact dates and actual availability remain unverified. Nominate and confirm before buyer submission.
- Final fixed-price deliverables, risk allowance and the proposed non-DE rates require offer-owner review. This is not a human approval or a request to send.