Varma Mutual Pension Insurance Company

Three models in production on one reusable MLOps pattern, built with your new specialists

Draft staffing revision - candidate evidence, scope and rates require owner confirmation

Three MLOps use cases over six months, with a funded three-role team

- **One MLOps golden path** on your existing cloud data platform replaces isolated deployment scripts and per-model configurations
- **Quick win:** use case 1 in production with monitoring by the end of month 2
- **Your new specialists own it:** they co-build use case 2 and build use case 3 themselves, with us reviewing
- **Three complementary roles, subject to candidate verification;** your right to require replacement accepted
- **Draft EUR 199,350 capacity fee**, EUR 66,450 per two-month milestone; scope, data readiness and risk allowance require confirmation.
- Leadership quick win: lead time from approved model to production, measured from use case 1 and reported monthly

One MLOps golden path for every model, instead of one-off deployment work per experiment

- Every model, current and future, uses the same repo template, configuration schema, pipeline and monitoring. A new use case starts from the template, not from a blank script.
- Built on your existing cloud data platform; MLflow-compatible registry, so the pattern is not tied to one vendor.

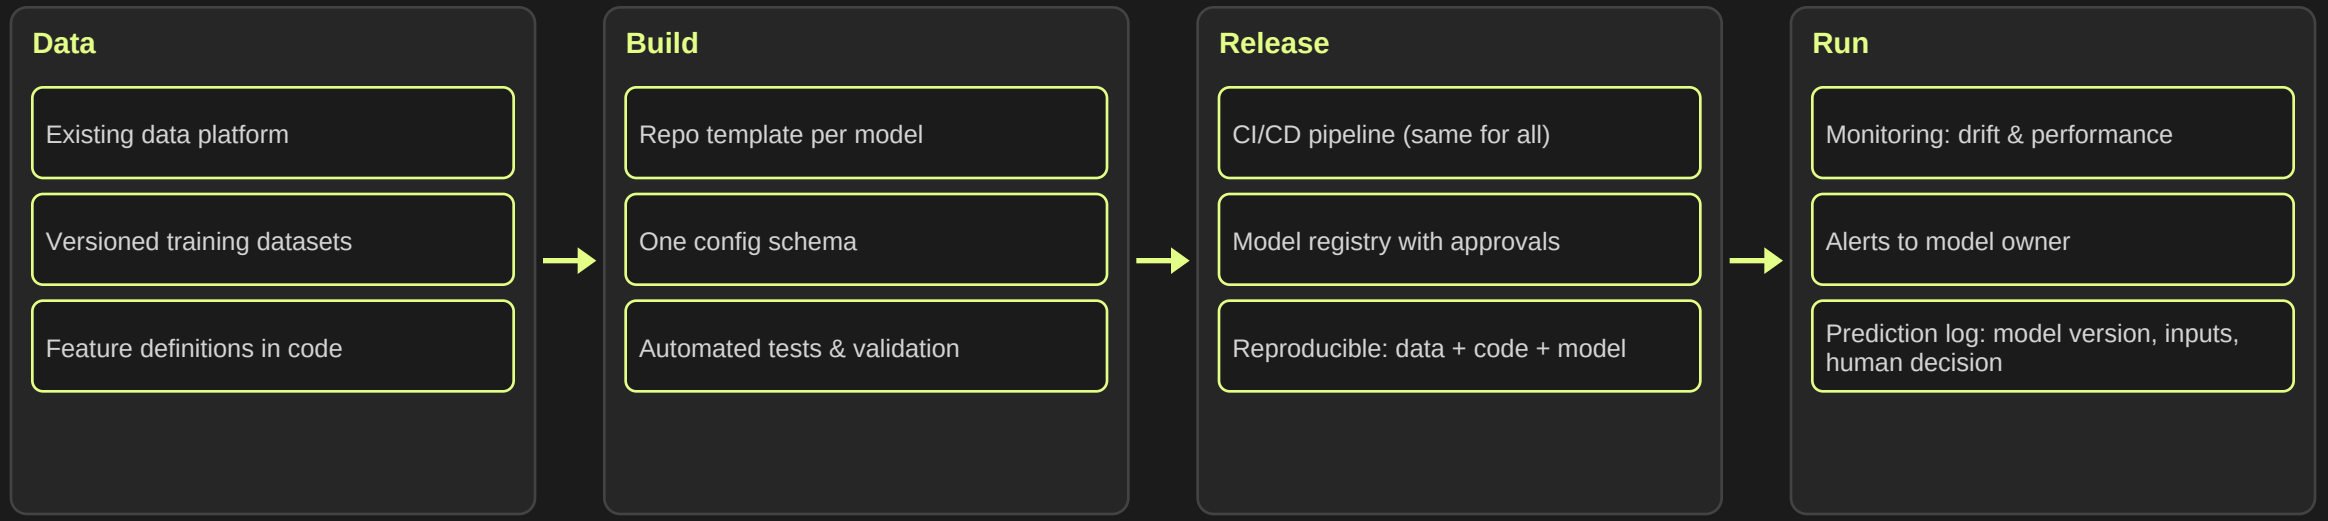

Models support pension decisions, they do not make them: shadow run first, then a human-approved release; rollback to the previous registered model in one step

Use case 1 in production by month 2; your new specialists ship use case 3 themselves

Month	Use case	Who builds	Technical proof
1–2	Use case 1 on the golden path (quick win)	Recordly, your specialists pairing	Live with drift alerts; performance checked when outcomes arrive (labels lag)
3–4	Use case 2 on the same template	Your specialists, Recordly pairing	Built without new deployment scripts
5–6	Use case 3 + triage of the other ~17 experiments by data readiness, value and risk class	Your specialists lead, Recordly reviews	Handover check: your team releases unaided

Gate 1 (week 3): use case 1 has usable data, a risk class and a named model owner; otherwise we swap to the next use case. Technical unknowns we price in: which of the three use cases has production-ready data; access to personal data for training (minimised, pseudonymised); EU AI Act classification per use case. Documentation uses Finnish pension terms (työeläke, TyEL, eläkepäätös), no translation needed.

Proposed role mix, allocations and verification before commitment

- **Replacement:** at your request, equal seniority within 10 working days, 2-week overlap at our cost
- **Delivery governance:** monthly steering with a written status (use cases, model health, risks); no weekly status meetings
- **Handover:** runbooks, the template and onboarding guide in Finnish, in your repos
- You may recruit our consultants after 12 months, with no fee

Proposed role	Role	Allocation	Required certificates	Target start
Lead ML engineer	Lead ML engineer	100%	Databricks Certified Machine Learning Professional; Microsoft Azure Data Scientist Associate	At signing
MLOps / platform engineer	MLOps / platform engineer	60%	Databricks Data Engineer Professional; Microsoft DevOps Engineer Expert	At signing
Solution architect	Solution architect, delivery governance	15%	Microsoft Azure Solutions Architect Expert	At signing

Draft role profiles: nominate consultants and verify CVs, certificates, languages, dates and actual availability before submission. No person is booked by this draft.

EUR 199,350 with all three AI delivery roles priced

- 20 working days/planning month; ML 100%, platform DE 60%, architect 15%. Confirm exact calendar and availability.
- Proposed ML and architect base EUR 125/h + EUR 5/h AI = EUR 130/h; DE EUR 115/h + EUR 5/h AI = EUR 120/h. Non-DE bases require owner confirmation.
- Three use cases prioritised at discovery; replace blocked use cases within the capacity limit. Scope growth and risk allowance require written agreement.

Role	Six-month effort / rate	Fee
Lead ML engineer	120 days x 7.5 h x EUR 130	EUR 117,000
MLOps / platform DE	72 days x 7.5 h x EUR 120	EUR 64,800
AI solution architect	18 days x 7.5 h x EUR 130	EUR 17,550
Total excl. VAT	210 days; 3 x EUR 66,450	EUR 199,350

No new licences assumed, subject to client confirmation. Candidate credentials and availability pending; VAT, cloud and travel excluded. No promise to productionise all 20 experiments.

Staffing review and decision before submission

- Rauno / Codex staffing review: maximum two architects respected; complementary roles included; DE/DA floors and AI premium applied to the stated role scope.
- This revision replaces the lower draft price for the same capacity. If the buyer budget is lower, reduce scoped work or hours and re-estimate; do not discount below hourly minima.
- Candidate CVs, credentials, language skills, exact dates and actual availability remain unverified. Nominate and confirm before buyer submission.
- Final fixed-price deliverables, risk allowance and the proposed non-DE rates require offer-owner review. This is not a human approval or a request to send.