RECORDLY

ERP pipelines automated into Data Vault 2.0, with a person approving every change

Priced per ingestion flow: you pay only for flows you accept, capped at €32,520 (excl. VAT). A small senior team plus Recordly's Virtual Data Engineer. No contractual deviations.

Proposal to Samuli Vainio, Head of Data Platform · Recordly Oy · Helsinki

Solution: your ERP into your Snowflake, built on what you already run

PART	HOW
Flow	ERP → your existing loader → Snowflake Raw Vault on AWS (hubs, links, satellites on your conventions) → your BI tool
Who builds	Boilerplate (loaders, satellites, tests) is drafted by the **Virtual Data Engineer** inside **your Snowflake account** on your AI credits. Our architect and engineer set the pattern and review.
Build or buy	**Your tools first.** We reuse your loader and dbt setup. We add nothing new unless a week-2 check shows a gap, and then you choose.
Out of scope	New platforms, new BI, ML models, and anything touching rail operations systems

> Real result from a Recordly project: a semi-complex pipeline built, tested and validated in **30 min instead of 12 h**.

What we read between the lines, and our safeguard

YOUR HISTORY	OUR SAFEGUARD
Last winter: an unchecked script wrote to production schemas	The agent **opens pull requests only**. It cannot approve or merge its own work. **A VR person approves every change.** Review agents test each PR, an auditor agent keeps the audit trail.
Hundreds of billed hours for boilerplate	**No hourly billing.** The agent writes the boilerplate; you pay a fixed price **per accepted ingestion flow**
Morning fleet reports must not break	Nothing is merged in the morning reporting peak. We keep the build lean, not feature-rich. Peak support = fast diagnosis, not new features
Alerts reach the right people	On a failed job the agent finds the root cause and **drafts a message** for fleet reporting users. A person reads and sends it. Nothing is sent automatically.
Diagnostic bots	Read-only on failed data jobs; it opens PRs, never touches rail systems

Delivery plan: flow by flow, each flow accepted and paid separately

MONTH	WORK	MILESTONE YOU ACCEPT
1	Pick 2-3 ERP pipelines with you. Measure baseline (lead time, review load). Deploy the VDE in your Snowflake. Architect sets conventions and 'perfect examples'.	**M1:** baseline + first pipeline through PR review
2	VDE drafts the remaining pipelines. Engineer reviews, tunes and tests with your engineers.	**M2:** agreed pipelines in dev, tests green
3	Your approvals, promotion on your schedule, Finnish runbooks, alert-to-draft flow	**M3:** pipelines live, runbooks done

Unknown: does your chat tool fit the alert thread? We check in week 1. If not, drafts go to your ticket queue instead. Knowledge transfer from week 1: pairing and all code in your repositories.

Team: available at signature, on site in Helsinki

ROLE	CERTIFICATIONS	LANGUAGES	DAYS
Jari Lehtonen, senior analytics engineer	SnowPro Advanced: Data Engineer · dbt Analytics Engineering · AWS Data Engineer Associate	fi (native), en	10
Tuomas Rantanen, solution architect	Azure Solutions Architect Expert	fi, en	8
Virtual Data Engineer (data engineer, review and auditor agents)	Runs 24/7 in your Snowflake	n/a	3 months

- **Majority of the team available from contract signature.** On site in Helsinki at least 2 days a week; all human days are on site.
- You may require replacement of any consultant. We replace at equal seniority, overlap at our cost.
- Documentation and walkthroughs in Finnish.

Price and terms: one total, no hourly billing

ITEM	BASIS	AMOUNT
Setup: conventions, DV 2.0 patterns, first flow as the "perfect example" (architect 4 d + engineer 2 d × €920)	one-off	€5,520
Each further ingestion flow, built by the Virtual Data Engineer, reviewed and approved by your engineers	**€2,700 per accepted flow** (up to 10 flows)	up to €27,000
Cap		**€32,520 excl. VAT**

- **Pay per flow, not per lump sum:** each flow is invoiced only when you accept it in production. You see the economics flow by flow, and you can stop after any flow with no further cost.
- The agent runs in your Snowflake on your AI credits; Snowflake credits per flow are reported, so you can compare cost per flow before ordering the next one.
- Snowflake AI credits are billed to you by Snowflake, not by us.
- Offer stands whether you award directly or run competitive tendering; we answer on your terms.
- **No contractual deviations.** Your standard terms apply as they are.