Elisa Oyj

Tech & architecture slides for the Elisa deck

Two slides for deck-writer: the architecture slide replaces the tier table on the solution slide, the gates slide goes right after it

Real-time where a domain needs it, on your own Google Cloud stack. Literal zero-batch is not achievable for reconciled finance

- **Freshness per domain, not one rule:** each domain gets the slowest tier that still meets its tested need. Target ranges per tier are agreed at Gate 1, after we measure current freshness.
- **Why not literally everything real-time:** finance close reconciles daily, and streaming all data raises run cost, failure surface and the retention exposure of traffic data. Option C, all real-time as written, is scored with its cost so the committee decides.
- **Existing services first:** Pub/Sub and change capture are proposed Google Cloud options, not assumed to be in your stack today. No new licence, no AI or ML, no new platform unless a tested need fails.

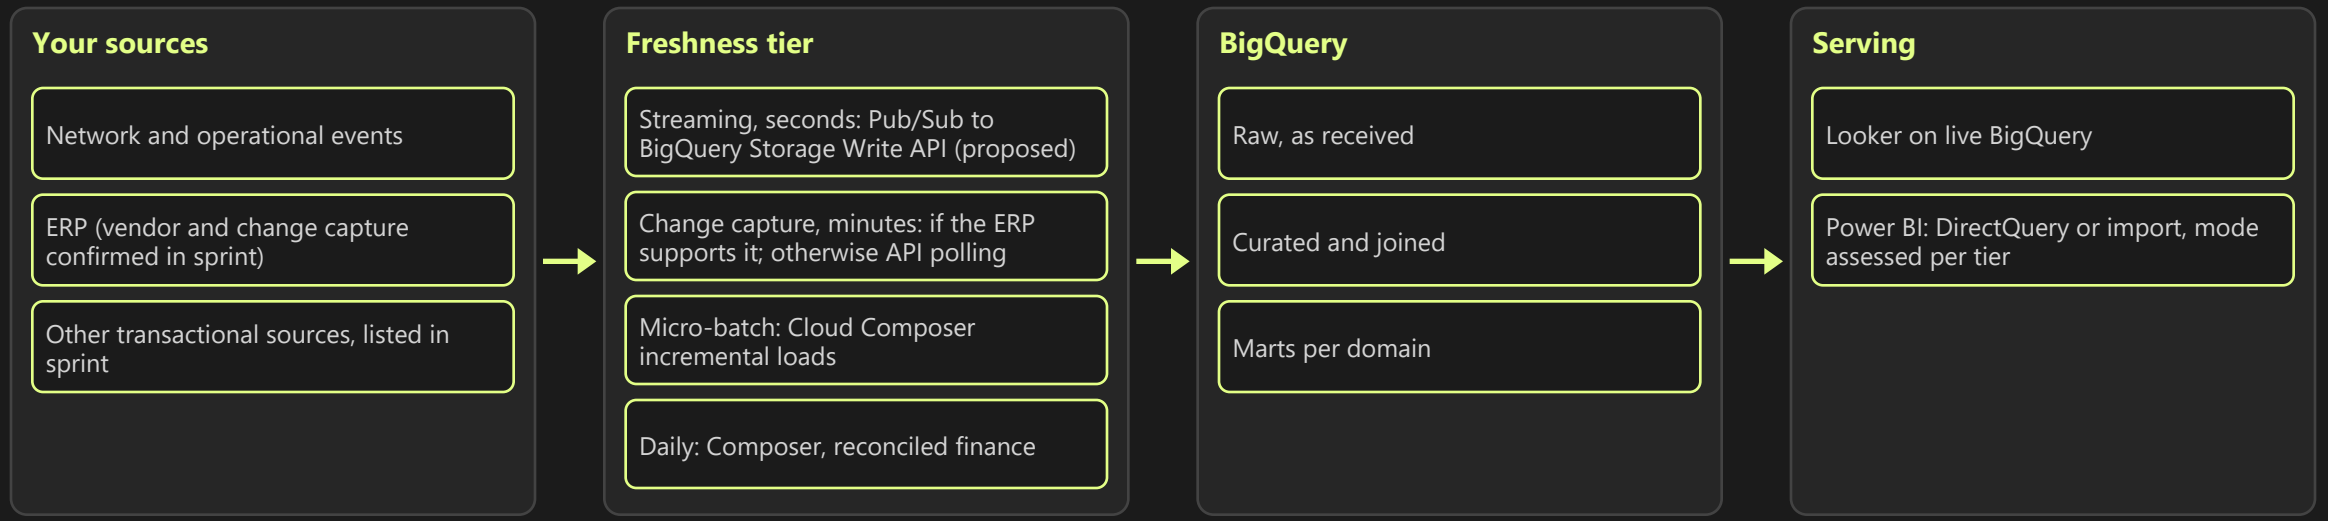

DAGs and CI/CD in your repos · same IAM and policy tags in Looker and Power BI · cost per tier · we read metadata only; no subscriber rows leave your GCP project

Technical gates: nothing is locked before it is measured

Gate	When	Passes when
Access ready	Before sprint day 1	NDA and secrecy undertaking signed; read-only Elisa-issued roles on BigQuery, Composer and Looker metadata and logs; no subscriber rows leave your GCP project; a sandbox dataset for the latency test
Gate 1	End of week 2	Each priority domain (3 to 5) has an owner, a stated freshness need and current freshness measured from logs; one real feed run Pub/Sub to BigQuery to Looker, end-to-end latency measured; ERP named with its change-capture options; target range per tier agreed
Checkpoint B	Week 8	Options A tuned, B tiered, C all real-time scored on latency, cost and quota per tier, Composer load, failure mode and effort; every score cites a measurement or is marked 'not tested'
Final	Week 10	Target architecture and roadmap accepted by the steering committee; anything untested marked 'not tested'

Unknowns and how each is settled: which ERP and its change capture (sprint; without CDC, API polling or micro-batch at an interval set at Gate 1) · Power BI connectivity (import refreshes on a schedule; DirectQuery freshness is bounded by query load, limits and cost; measured per report in the sprint) · real freshness need per domain (owners answer in your chat tool by week 2) · where Composer code, CI/CD and access policies live (recorded in the sprint; existing pipelines stay unless a test fails).