VR-Yhtymä Oyj (VR Group)

Tech & architecture slides for the VR Group deck

Two slides for deck-writer to insert after the answer-first slide

Every source gets a named extraction method, landing in your Data Vault without new tools

- We propose the extraction for all sources, as asked: one metadata-driven loading pattern on AWS, configured per source, not one script per feed.
- Business keys follow the railway terms your operations people already use, documented in Finnish.

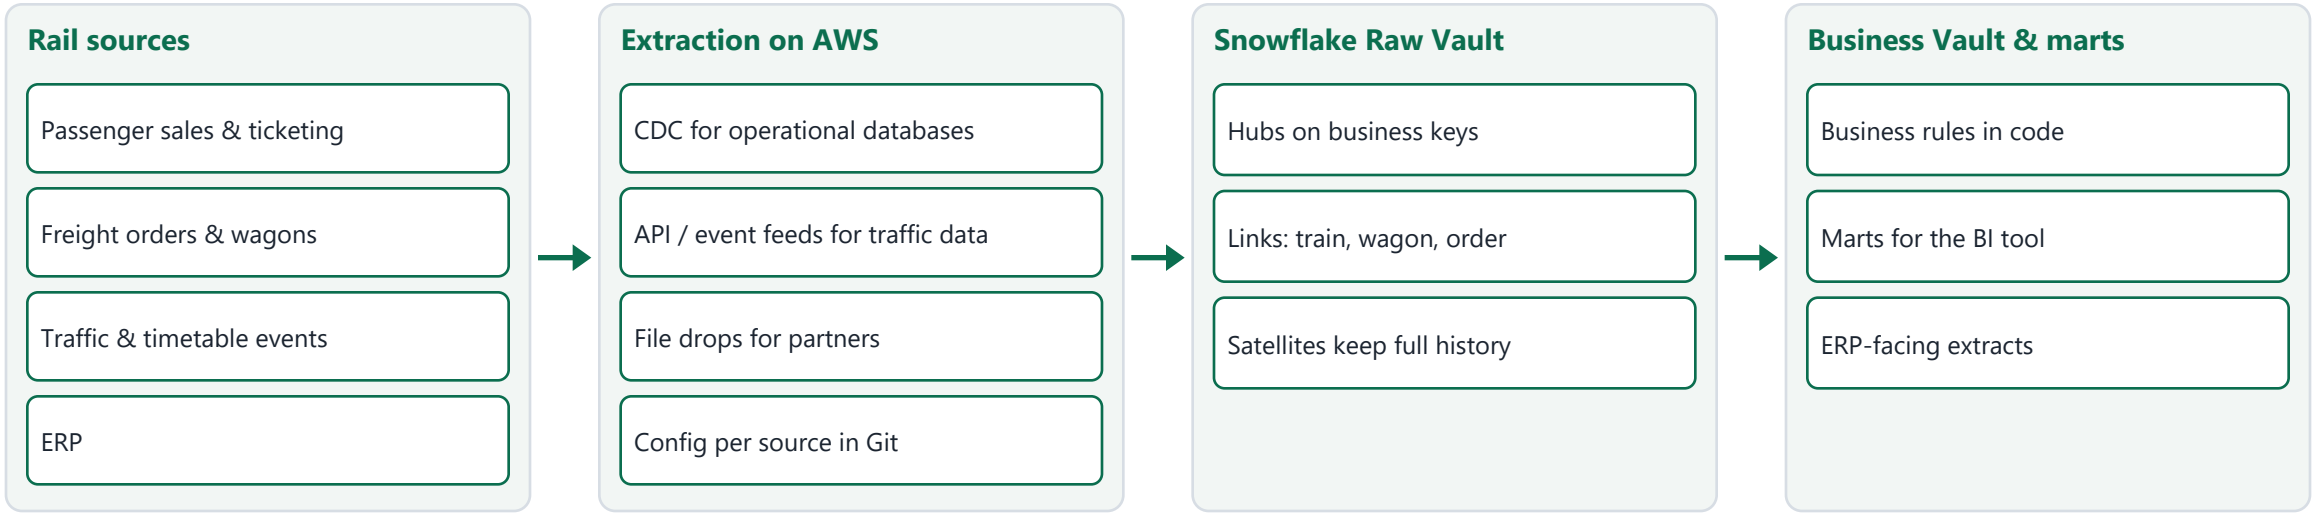

Q1: build and test in non-production only. Production releases in December and from April, each with a rollback

Technical gates protect operations: nothing reaches production in Q1

Gate	When	Pass criterion
Extraction inventory	Week 2	Every source has a method, owner and load frequency
First satellites live	Before 20 Dec	Row counts reconcile to source, no load errors for 5 days
Q1 release package	End of March	All Q1 work tested in non-production, release plan approved
April release	Agreed weekend window	Reconciliation green; rollback tested

Technical unknowns: source system access on AWS, change-data-capture support per database, the current Data Vault loader (we use yours if one exists). Snowflake makes Recordly's Virtual Data Engineer a possible add-on (see team/vde).